



UNIVERSITY OF PISA

MSc in Computer Engineering

Project for Performance Evaluation of Computer
Systems and Networks

Multi-core scheduling

Professors:

Prof. Giovanni Stea

Ing. Giovanni Nardini

Students:

Taulant Arapi (645308)

Francesco Barcherini (645413)

Antonio Ciociola (645324)

ACADEMIC YEAR 2024/2025

Contents

1	Photos	4
2	Introduction	5
2.1	Description and specification	5
2.2	Objectives	6
2.3	Indices	6
2.4	Assumptions?	7
2.5	Preliminary stability estimation	7
3	Implementation	9
4	Verification	11
4.1	Code verification	11
4.2	Degeneracy test	11
4.3	Consistency test	12
4.4	Continuity test	13
5	Analysis	14
5.1	Calibration of the parameters	14
5.2	Time parameters setup	15
5.2.1	Warm-up period	15
5.2.2	Simulation duration	16
5.3	Subsampling	17
5.4	Statistical analysis	18
5.4.1	First-Come-First-Served	18
	Turnaround time	18
	Response time	21
	Throughput	21
	CPU utilization	21
	Queue length	21
5.4.2	Shortest-Job-First	21
	Turnaround time	21
	Response time	21

CONTENTS	3
Throughput	21
CPU utilization	21
Queue length	22
6 Conclusions	23

1 Photos

SCADENZA: 21/01/2025

TEST

- objective
- indices
- descriptions
- assumptions
 - infinite queue
- validation (assumptions)

Factors

- mean generation time (us) → 1, 2, 5, 10
- mean process duration → 1, 5, 25, 125
- number of processes → 1, 2, 4, 8, 16
- ϕ -cutoff → 0, 0.5, 0.5, 0.75, 1
- isrcfg = true/false
- generation_type = constant/uniform/exponential
- duration_type = constant/heavy-tailed/exponential

Verification

- debugging, EV
- test-case facility
- consistency test (mcpus, times...)
- degeneracy test
- continuity test

Calibration of parameters DATA AVAILABLE

duration → tanto

Warm-up → coda con $f(N, P)$ processi
togliere la prima porta
analizzando mean e std

mel codice ≠ Seeds, same scenario
avg window

turnaround time

- waiting time
- utilization CPU
- CPU active mel tempo (t)
- Processi in coda pronti (t)

AVE

Dopo Warm-up grafici

isrcfg = true/false

- generation_type = constant/uniform/exponential
- duration_type = constant/heavy-tailed/exponential

Verification

- debugging, EV
- test-case facility
- consistency test (mcpus, times...)
- degeneracy test
- continuity test

Calibration of parameters DATA AVAILABLE

duration → tanto

Warm-up → coda con $f(N, P)$ processi
togliere la prima porta
analizzando mean e std

mel codice ≠ Seeds, same scenario
avg window

turnaround time

- waiting time
- utilization CPU
- CPU active mel tempo (t)
- Processi in coda pronti (t)

AVE

Dopo Warm-up grafici

2 Introduction

2.1 Description and specification

It is required to analyze a multi-core computer equipped with N CPUs, which execute multiple interactive processes. The processes are dynamically generated at intervals of T seconds. Each process has a total duration D , which is divided into three distinct execution phases:

1. initial processing phase;
2. I/O operation phase;
3. final processing phase.

The times T and D are defined as IID random variables. The processes are categorized into two types and generated as *CPU bound* with probability p or *I/O bound* with probability $1 - p$. The difference is that:

- in CPU-bound processes the I/O operation phase constitutes 20% of the total duration D , and the other two phases both the 40%;
- in I/O-bound processes the I/O operation phase constitutes 80% of their total duration D , and the other two phases both the 10%;

Processes are assigned to CPUs by the operating system's scheduler, which selects tasks from a list of "ready" processes. When a CPU becomes idle for entering the I/O operation phase or ending the final processing phase of a process, the scheduler immediately assigns it to a new ready process, if present. Once a process completes its I/O operation phase, it is marked as ready again and is returned to the scheduling queue. A process exits the system once it has successfully completed all three phases of execution.

The scheduler can be implemented to follow two distinct policies:

- First Come First Served (FCFS): processes are scheduled in the order of their arrival after the generation or after the end of the I/O phase;
- Shortest Job First (SJF): scheduling is determined based on the shortest remaining time until either the process's I/O phase or its completion.

2.2 Objectives

The aim of this project is to assess and analyze the execution of a multi-core scheduling system under varying operational scenarios. The objective is to determine the optimal characteristics of the systems to make it work without the need to discard processes. Furthermore, the analysis aims to assess the conditions under which the best system behavior is achieved in terms of resource use efficiency and time performance evaluation. In particular, the different configurations taken into consideration are the scheduling policies, the number of CPUs in the system, the type of the processes (CPU bound or I/O bound), the generation intervals and the duration of processes.

The overarching goal is to generate insights into optimal scheduling strategies and resource allocation practices for multi-core environments, enhancing system throughput and minimizing delays.

2.3 Indices

To evaluate the system's performance comprehensively, the following indices will be analyzed in detail:

- turnaround time R : defined as the total time elapsed from the arrival of a process in the system to its completion. This metric provides a measure of the computer's overall responsiveness and of the time required for different types of processes to be executed;
- waiting time W : the cumulative time [!:: COME LO CALCOLIAMO?] a process spends in the ready queue before being assigned to a CPU for execution. This metric represents the impact on the turnaround time of the queuing policy and of the workload of the system;
- CPU utilization of the n_{th} CPU ρ_n : the fraction of time that each CPU is actively engaged in executing processes instead of being idle. This metric reflects the efficiency of resource usage within the system;
- active CPUs over time N_A : a dynamic measure of the number of CPUs actively processing tasks at any given moment. This index helps assess load distribution and system evolution;
- queue length over time N_q : tracks the size of the ready queue throughout the simulation, highlighting bottlenecks and variations in process scheduling.

These indices will be statistically analyzed to identify trends, anomalies, edge cases and key factors influencing system performance. By examining these metrics under diverse configurations of N , p and scheduler policies, the project aims to provide actionable recommendations for improving the efficiency and adaptability of multi-

core scheduling systems.

2.4 Assumptions?

To perform the analysis, the following assumptions are made:

- There is sufficient memory to store the list of processes ready for execution, so the ready queue will be considered as infinite.
- The scheduler itself has a negligible load on the CPUs; there is no delay between the executions of the modeled processes. [!!: assumiamo che sia zero? if not, quanto facciamo?]
- With the shortest job first (SJF) scheduling, an accurate estimate of the execution times is known a priori. If the scheduler is FCFS, this data is not used.

2.5 Preliminary stability estimation

The system can be modeled as a queueing network as shown in Figure 2.1.

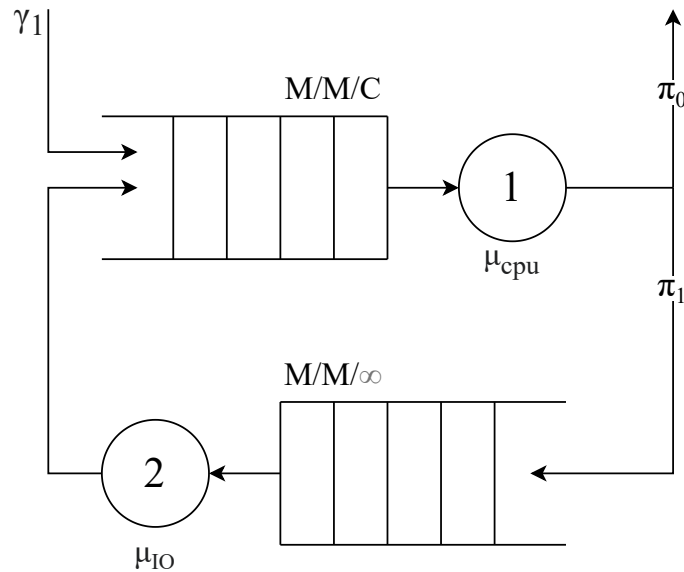


Figure 2.1: Model of the system as a queueing network

Let's analyze the system when the process generation time T and the process duration D are exponentially distributed, with mean values $E[T]$ and $E[D]$ respectively, and when the queues follow a FCFS scheduling policy. New processes are generated at rate $\gamma_1 = \frac{1}{E[T]}$. The initial processing phase and the final processing phase are modeled by the $M/M/C$ Service Center 1, where $C = N$ is the number of CPUs. Here it is assumed that the service time t_{cpu} can be $0.4D$ for CPU bound jobs and $0.1D$ for I/O bound jobs. This means that, for total probability theorem, the mean

value of the service time of the CPU processing is

$$E[t_{cpu}] = 0.4E[D] \cdot p + 0.1E[D] \cdot (1 - p) = 0.3E[D] \cdot p + 0.1E[D]$$

so that

$$\mu_{cpu} = \frac{1}{E[t_{cpu}]} = \frac{1}{0.3E[D] \cdot p + 0.1E[D]}$$

After the CPU processing phase, we can assume that it is equally likely that the process is at the beginning of the I/O phase or at the end of final processing phase. In the second case, it leaves the system with routing probability $\pi_0 = \frac{1}{2}$. If the process starts the I/O phase, it is routed with probability $\pi_1 = \frac{1}{2}$ to the $M/M/\infty$ Delay Center 2 with service time t_{IO} . Reasoning as before, the mean value of the service time of the I/O processing is $0.2D$ for CPU bound jobs and $0.8D$ for I/O bound jobs. The mean value of the service time of the I/O processing phase is

$$E[t_{IO}] = 0.2E[D] \cdot p + 0.8E[D] \cdot (1 - p) = 0.8E[D] - 0.6E[D] \cdot p$$

so that

$$\mu_{IO} = \frac{1}{E[t_{IO}]} = \frac{1}{0.8E[D] - 0.6E[D] \cdot p}$$

Both service centers are of type $M/M/C$, the external arrival γ_1 is Poissonian, routing probabilities are state-independent and arcs are traversed in zero time. Therefore the hypothesis of Jackson's theorem are satisfied and the system admits a product form if the stability conditions $\rho_i = \frac{\lambda_i}{C_i \cdot \mu_i} < 1$ are satisfied for each service center i . To compute λ_1 and λ_2 we can define the routing matrix $\Pi = \begin{bmatrix} 0 & \frac{1}{2} \\ 1 & 0 \end{bmatrix}$. Arrivals are $\gamma = \begin{bmatrix} \gamma_1 \\ 0 \end{bmatrix}$. It is $\lambda = (I - \Pi^T)^{-1} \cdot \gamma$. We can observe that, since the input and output must balance, $\gamma_1 = \frac{\lambda_1}{2}$ and from the routing we get $\lambda_2 = \frac{\lambda_1}{2}$. At the end we get $\lambda_2 = \gamma_1 = \frac{1}{E[T]}$ and $\lambda_1 = \frac{2}{E[T]}$.

The stability condition for the delay center is always true, so the system is stable if the first service center is stable. To have $\rho_1 < 1$ we get $\frac{\lambda_1}{N \cdot \mu_{cpu}} < 1$: substituting and rearranging we get

$$N > \frac{3p + 1}{5} \cdot \frac{E[D]}{E[T]} \quad (2.1)$$

The metrics to be computed can be obtained from the analysis of the two service centers. The mean number of busy CPUs is, for example,

$$E[N_{cpu}] = \frac{\lambda_1}{\mu_{cpu}} = \frac{3p + 1}{5} \cdot \frac{E[D]}{E[T]} \quad (2.2)$$

And, with some computation based on known formulas, it is possible to get a theoretical estimation of the mean number of processes in the ready queue $E[N_q] = E[N_{q1}]$, of the mean waiting time in the ready queue $E[W] = 2E[W_1]$ and of the turnaround time $E[R] = E[W] + E[D]$.

3 Implementation

We implemented the Computer system in OMNeT++.

Figure [Figure 3.1](#) provides an overview of the system implementation. Below is a detailed explanation of each module:

- **ProcessGenerator**: this module generates processes according to a specified interarrival time distribution. Each process is assigned a total duration D , and the duration of each phase (initial, I/O, and final) is set to $0.4D$, $0.2D$, and $0.4D$ if the process is CPU bound, or $0.1D$, $0.8D$, and $0.1D$ if the process is I/O bound. The process is then sent to the scheduler.

- **scheduler**: this module contains an infinite-capacity queue. If `isFCFS` is set to `true`, a FIFO queue is used. Otherwise, a priority queue is used, with processes ordered by increasing duration until the I/O phase or termination (shortest job first).

When a new process arrives, the scheduler sends it to the first free CPU. If no CPU is free, the process is enqueued. When a CPU is freed, the next process in the queue is sent.

When a process enters the I/O phase, the scheduler waits for the specified I/O phase duration and adds the process back into the queue.

When a process terminates, its turnaround time and waiting time are emitted as signals.

The scheduler also keeps track of the number of active CPUs and the queue length, emitting those every time a message is received.

- **cpu**: There are N CPUs which operate independently. Each one, when it is free, receives a process from the scheduler and simulates its execution. When it finishes the execution phase (i.e. the I/O phase is reached or the process terminates), a message is sent to the scheduler, indicating that the CPU is now free. Every CPU records its utilization, emitting the busy time as a signal when the CPU is freed.




Figure 3.1: Overview of the system implementation.

4 Verification

4.1 Code verification

The code has been verified to work as expected with the debug GUI that comes with OMNeT++. In particular, the code has been tested trying to execute all the possible paths of the code in order to simulate all different cases, and comparing the actual state of the system with the expected one.

For the analysis, the specific constructs of OMNeT++ (such as `cQueue` for the queues or user defined messages) allowed to better follow the flow of events in the simulations. In fact the inspection of the content of messages and the state of the queues has been crucial to understand the behavior of the system.

Additional information has been gathered printing log messages in the console with the use of `EV`, in order to have a better comprehension of the flow of the simulation. For example, `EV` has been very useful to track the main events of the simulation, such as the arrival of a new process, the start and end of the three phases of execution, the assignment of a process to a CPU. A simple graphical representation of the processes queued waiting to be assigned to a CPU helped to verify the correct behavior of the scheduler in both FCFS and SJF policies.

4.2 Degeneracy test

Degeneracy tests aim to verify the correct behavior of the system at extreme values of the parameters. The main quantitative parameters are the number of CPUs N , the probability of a process to be CPU bound p , the mean generation time $E[T]$ and the mean process duration $E[D]$. The possible scenarios can be tested with FCFS or SJF scheduling policy and with exponential or uniform distributions of the two times.

The probability p set at 0 or 1 leads to the generation of only I/O bound or CPU bound processes. The simulation works as expected in the two cases with different values of the statistics (ex. turnaround and waiting time, number of busy processors etc.), without affecting the functionality of the system.

When the number of CPUs is set to 0, no CPU is instantiated and the processes are never assigned to a CPU. The system works as expected, with the processes waiting

in the queue until the finish of the simulation time or the end of memory availability. Tests with only 1 CPU show a good behavior and no problem related to the handling of vectors of gates with length 1. The test with a huge number of CPUs (100) is not different from every test with more than one CPU with respect to the multicore functionalities of the computer; obviously, in this case, there is always an available CPU to assign a process to and there is no queueing.

Talking about the mean generation time $E[T]$ and the mean process duration $E[D]$, what matters is the ratio between the two. To simulate scenarios with very large (small) generation time and very small (large) process duration, the two parameters have been set to 10^3ms and 10^{-3}ms respectively, and vice versa. In the case of $E[T] \gg E[D]$, every process executes its three phases before the generation of a new one, so there is no queueing as expected. In the case of $E[T] \ll E[D]$, the processes are generated at a very high rate so it is necessary to set `sim-time-limit` to a small value to end the simulation in a reasonable time. The system works as expected, with a high number of processes in the queue and CPUs always busy.

4.3 Consistency test

The consistency test aims to observe a similar behavior of the system when the parameters are changed in such a way that the result is expected to be still the same.

With the hypothesis of First Come First Served scheduling policy and exponential distributions of times, the first tests verify the scenarios with fixed $p = 0.5$ and $N = 6$ and changing values of $E[T]$ and $E[D]$, keeping the ratio $\frac{E[D]}{E[T]}$ still constant and equal to 10. In this way, the generation rate and the service rate change in the same way and the statistics related to the utilization should remain the same. In particular, the values chosen are 1ms, 4ms and 10ms for $E[T]$ and 10ms, 40ms and 100ms for $E[D]$. The simulation effectively shows the same CPU utilization around the value of 0.835. Also the mean number of busy CPUs is constant and equal to 5 and the mean number of processes in the ready queue is around 4.7. The mean turnaround time, waiting time in the queue and service time show, instead, a proportionality with $E[D] \propto \frac{1}{\mu}$ as expected. The times are multiplied with a factor 4 and 10 in the second and third scenarios with respect to the first one.

A second continuity test consists in changing also the number of CPUs N , changing the ratio $\frac{E[D]}{E[T]}$ but keeping constant $\frac{E[D]}{E[T] \cdot N}$, proportional to the utilization of the CPUs $\rho = \frac{\lambda}{N \cdot \mu}$. The chosen values are 3, 6 and 9 for N , 1ms, 4ms and 10ms for $E[T]$ and 5ms, 40ms and 100ms for $E[D]$. The CPU utilization extracted from the simulation effectively remains the same in the different scenarios and equal to 0.835. The mean number of busy CPUs increases proportionally with $\frac{E[D]}{E[T]}$: 2.5, 5 and 7.5 in the three cases. The other metrics change without a clear proportionality with the parameters, still consistently with the changes and with the expectations for a M/M/C system.

Another important test consists in changing these parameters in conditions of

instability. For example, in the latter case with N equal to 1, 2 and 3 the system keeps being unstable: the CPUs cannot handle the load of the processes and the queue grows indefinitely.

The same tests have been performed combining the scenarios with the Shortest Job First queueing policy and with uniform distributions of the times. The results confirm the consistency of the system with the expected behavior in term of CPU utilization and stability conditions. Of course there is no more the same proportionality between the times and the metrics as regards queueing, while it is still valid for the service times of the processes.

4.4 Continuity test

The continuity test verifies that a slight change in the input does not cause wild changes on the system behaviour.

In the system, the test is handled in both FCFS and SJF scheduling policies. The parameters vary with small steps of 5% of their value, so that N takes values 19, 20 and 21, p takes values 0.475, 0.5 and 0.525, $E[T]$ takes values 1.425ms, 1.5ms and 1.575ms and $E[D]$ takes values 47.5ms, 50ms and 52.5ms. The system is stable in all the cases, as expected, and the variation of the metrics is around no more than 10% of their value. For example, the CPU utilization varies from a minimum of 0.817 to a maximum of 0.854 (4.33% of variation). The turnaround time goes from 49.7ms to 55.4ms (10.3%), and similarly do the other statistics.

5 Analysis

5.1 Calibration of the parameters

Having a lot of parameters to calibrate, we decided to limit the number of possible values for each parameter. This choice was made to reduce the number of simulations to be performed and to have a more manageable number of results to analyze. The parameters that we decided to calibrate are the following:

- **pCpuBound:** {0.1, 0.9}
The probability of a process being CPU-bound has two values, one for a system with a high number of I-O bound processes and the other for a system with a high number of CPU-bound processes.
- **meanProcessDuration:** {0.1 s}
We chose a fixed value for the mean process duration, by doing this we can focus on the effect of the other parameters on the metrics.
- **meanGenerationTime:**
The mean generation time was chosen to have a comparable load on the system for both values of **pCpuBound**. For each value of **pCpuBound** there is a simulation with heavy load and one with a lighter load. The values are in the order of tens of milliseconds.
- **numCpus:** {4, 12}
The number of CPUs was chosen to simulate quad-core, and newer systems with 12 cores.
- **isFCFS:** {true, false}
The scheduling policy that the scheduler uses can be either First-Come-First-Served or Shortest-Job-First.
- **generationType:** {"exponential", "uniform"}
In addition to the exponential distribution, we also considered the uniform distribution for the generation of processes.
- **durationType:** {"exponential", "uniform"}
In addition to the exponential distribution, we also considered the uniform distribution for the duration of the processes.

5.2 Time parameters setup

5.2.1 Warm-up period

To determine the warm-up period, we observed the development of the mean number of busy CPUs over time, through 10 independent runs and with different parameters configurations. The number of busy CPUs is a good indicator of the system's stability, as it doesn't depend on the scheduling policy. Our observations showed that the system stabilizes well before 200s in all tested configurations. To handle variability and ensure robustness in worse cases, we set the warm-up period to 200s.




Figure 5.1: Line chart of the mean number of busy CPUs over time for different configurations of meanGenerationTime and meanProcessDuration and multiple repetitions.

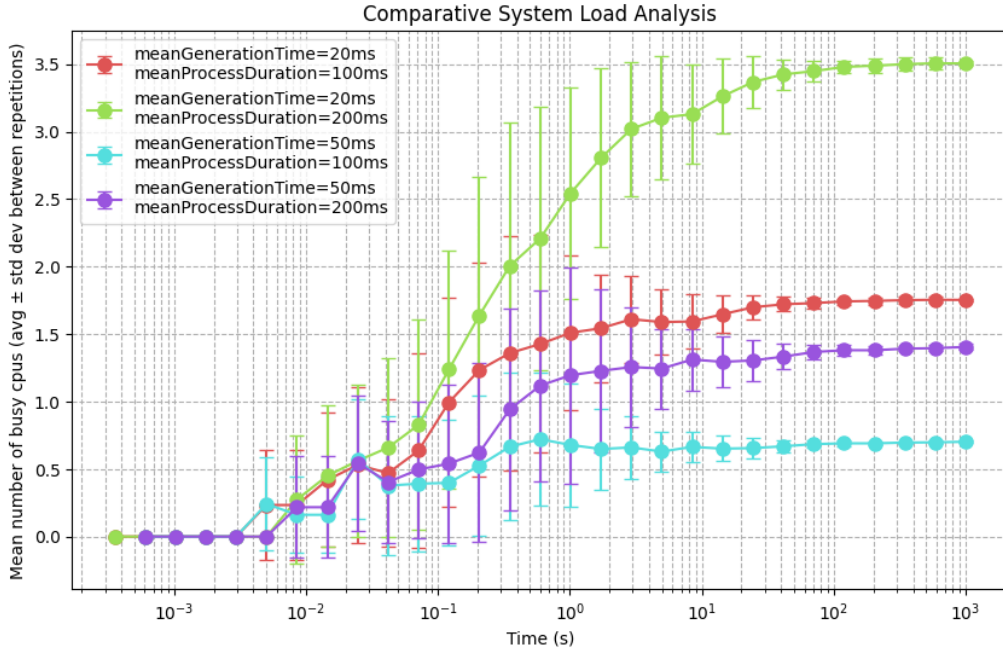


Figure 5.2: Average and Std dev of the mean number of busy CPUs over the repetitions.

Time (s)	20ms, 100ms		20ms, 200ms		50ms, 100ms		50ms, 200ms	
	Avg	Std Dev	Avg	Std Dev	Avg	Std Dev	Avg	Std Dev
1	1.52	0.58	2.54	0.78	0.68	0.46	1.20	0.80
2	1.54	0.37	2.85	0.64	0.67	0.27	1.25	0.51
5	1.59	0.24	3.10	0.46	0.63	0.15	1.24	0.29
10	1.63	0.18	3.18	0.34	0.65	0.10	1.29	0.20
20	1.68	0.11	3.33	0.22	0.64	0.08	1.27	0.15
50	1.72	0.05	3.43	0.10	0.68	0.05	1.35	0.09
100	1.74	0.02	3.47	0.05	0.69	0.02	1.37	0.04
200	1.75	0.02	3.49	0.05	0.69	0.02	1.38	0.03
500	1.75	0.03	3.50	0.05	0.70	0.01	1.40	0.02
1000	1.75	0.01	3.51	0.02	0.70	0.01	1.41	0.02

Table 5.1: Average and Std dev of the mean number of busy CPUs over the repetitions.

5.2.2 Simulation duration

For the simulation duration, we chose to end it at 1000s. This ensures that after discarding the initial 200s warm-up period, there remains 800s of simulation data for analysis. Given that the mean generation time is on the order of tenths of a second, this duration strikes a balance between obtaining statistically meaningful results even after subsampling and maintaining reasonable simulation times.

5.3 Subsampling

To analyze the data, the assumption of IID-ness will be needed, but without further modification the samples do not uphold it. For instance, if a process finishes with a large turnaround time, which is caused by a long queue, it is likely that the same thing will happen for the next processes.

To address this and ensure independence between samples, subsampling has been employed. The new sample is constructed taking each point of the starting sample with probability p . $p = \frac{1}{2^k}$ and k is the smallest integer that passes the Ljung-Box test. Using the Ljung-Box test makes it possible to automate the subsampling step. It was implemented to test the first 30 lags, with a significance level of 5%, with these values the sample is considered independent if:

$$Q = n(n+2) \sum_{k=1}^{30} \frac{\hat{\rho}_k^2}{n-k} < \chi_{0.95,30}^2 \approx 43.77 \quad (5.1)$$

As utilization increases, the correlation becomes stronger, as shown in fig. 5.3.

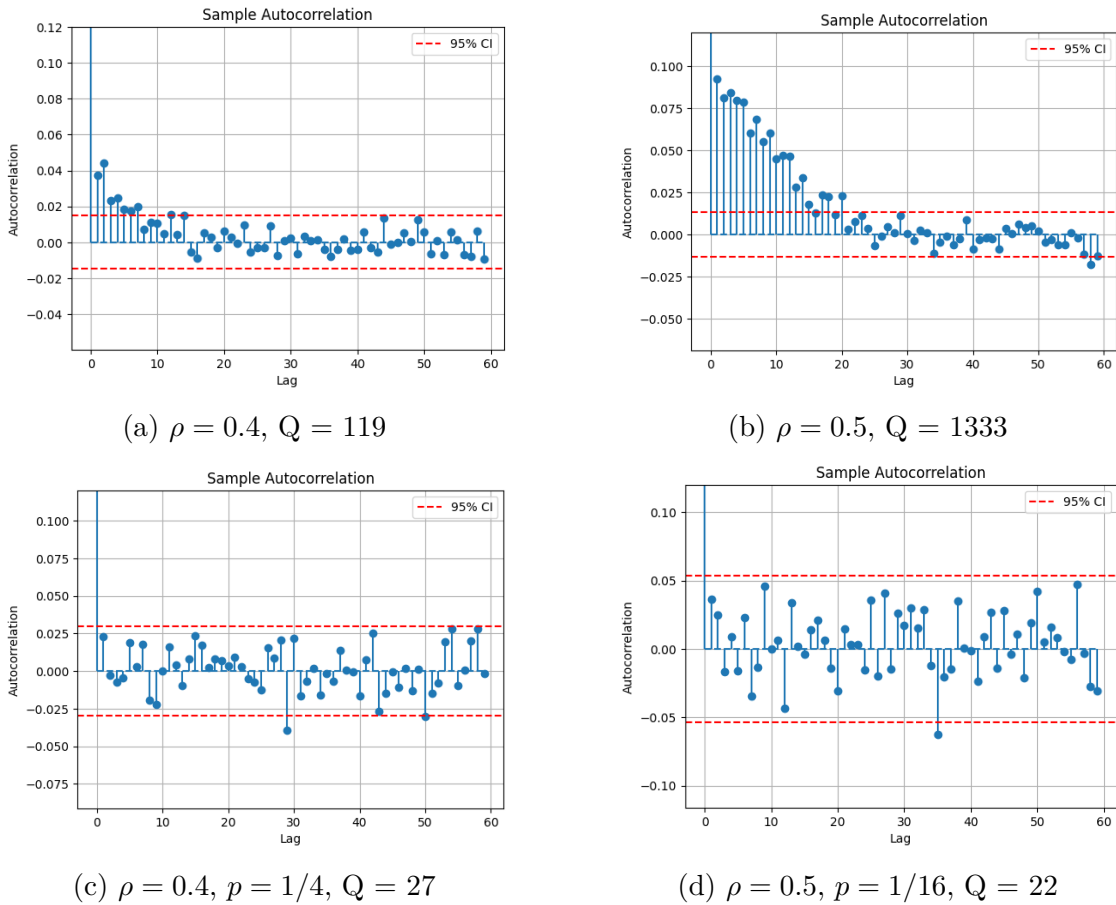


Figure 5.3: Comparison of turnaround time autocorrelation with and without subsampling for different loads.

5.4 Statistical analysis

5.4.1 First-Come-First-Served

Turnaround time




Figure 5.4: Empirical CDF of the turnaround time of the systems.




Figure 5.5: Density plot of the turnaround time of the systems.

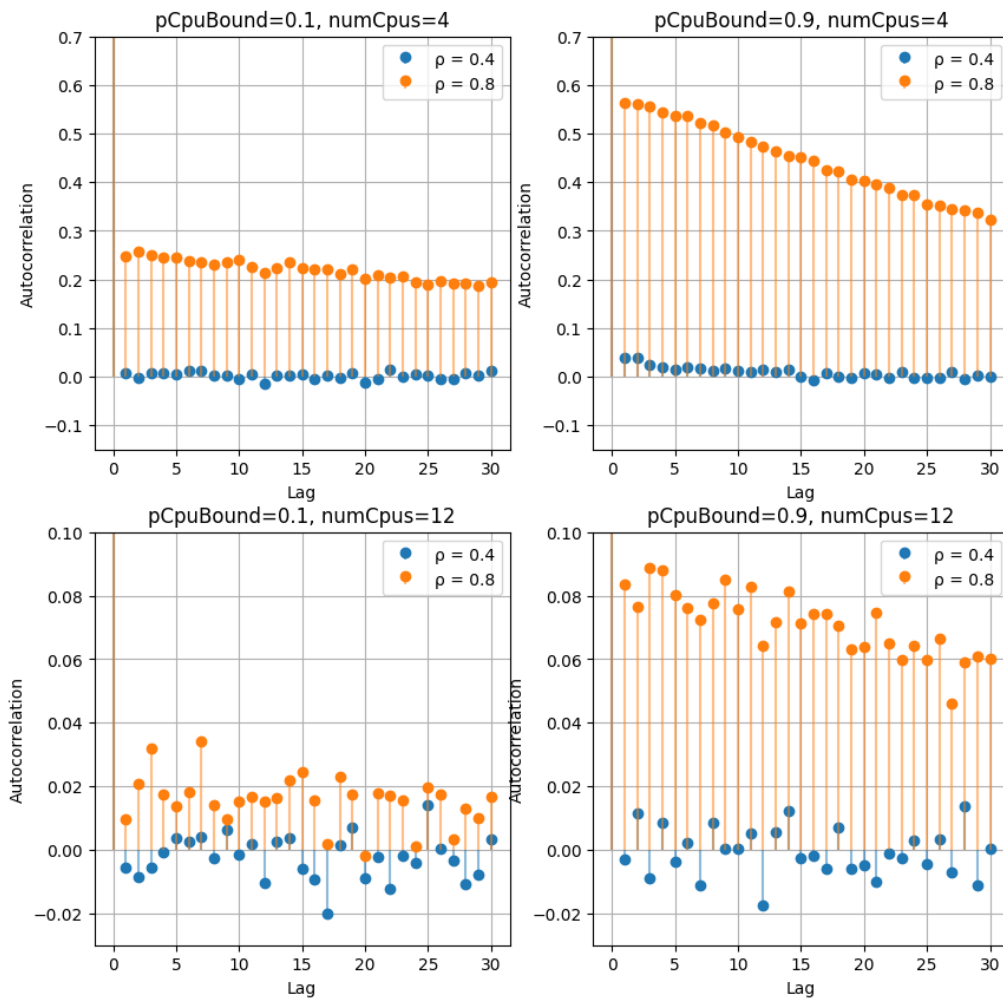


Figure 5.6: Autocorrelation plot of the turnaround time of the systems.

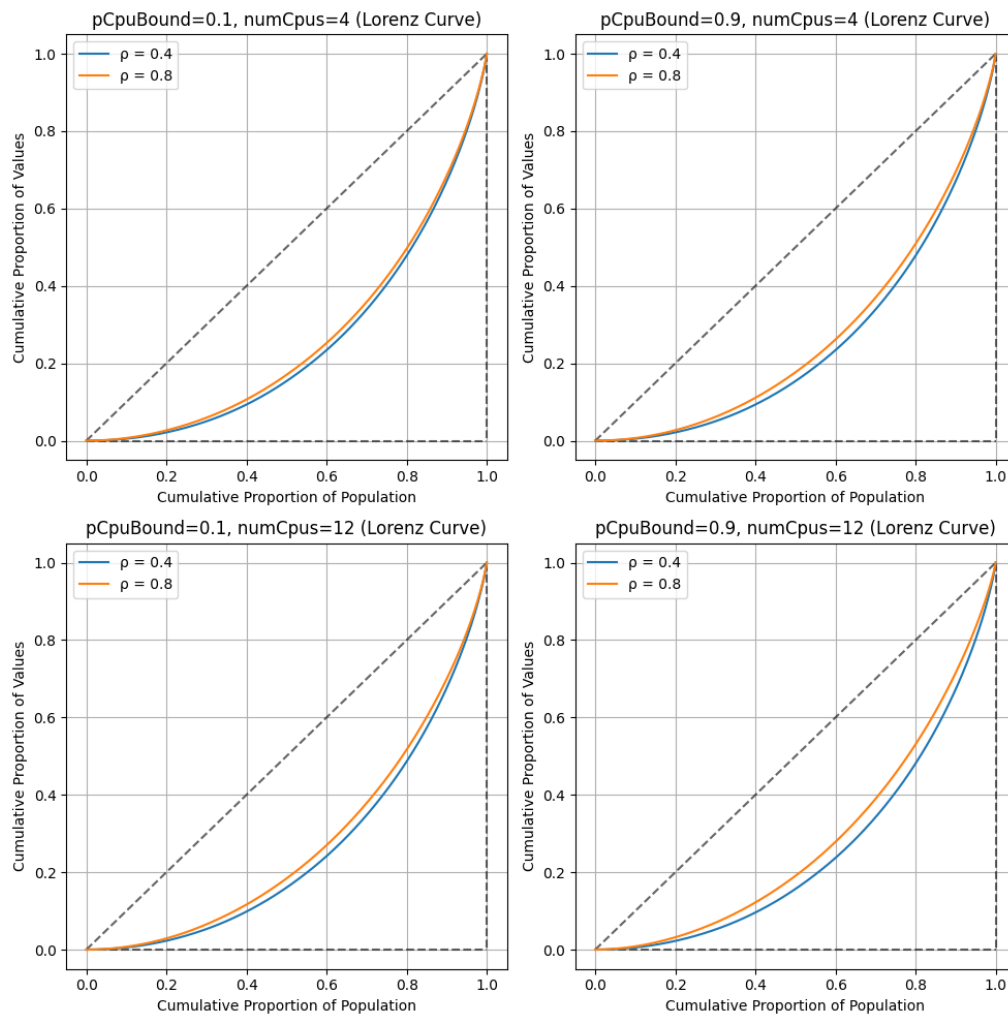


Figure 5.7: Lorenz curve of the turnaround time of the systems.

Response time

Throughput

CPU utilization

Queue length

5.4.2 Shortest-Job-First

Turnaround time

Response time

Throughput

CPU utilization

There is no difference with section 5.4.1 since the FCFS and SJF schedulers have the same behaviour in terms of CPU utilization.

Queue length

6 Conclusions

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

```
1 // Your First C++ Program
2
3 #include <iostream>
4
```

```
5 int main() {  
6     std::cout << "Hello World!";  
7     return 0;  
8 }
```

Listing 6.1: test.cpp

Nunc sed pede. Praesent vitae lectus. Praesent neque justo, vehicula eget, interdum id, facilisis et, nibh. Phasellus at purus et libero lacinia dictum. Fusce aliquet. Nulla eu ante placerat leo semper dictum. Mauris metus. Curabitur lobortis. Curabitur sollicitudin hendrerit nunc. Donec ultrices lacus id ipsum.

[Persona 1: Commento 1]

[Persona 2: Commento 2]

Test cite^[1].

Test footnote¹.



Figure 6.1: Un gattino.

Ciao

Nunc sed pede. Praesent vitae lectus. Praesent neque justo, vehicula eget, interdum id, facilisis et, nibh. Phasellus at purus et libero lacinia dictum. Fusce aliquet. Nulla eu ante placerat leo semper dictum. Mauris metus. Curabitur lobortis. Curabitur sollicitudin hendrerit nunc. Donec ultrices lacus id ipsum.

¹ Footnote

Bibliography

- [1] Google. *Google*. 2025. <https://google.com>.